

Advanced Programming

Erlang Introduction — The Sequel

Ken Friis Larsen
`kflarsen@diku.dk`

Department of Computer Science
University of Copenhagen

October 5, 2017

Today's Menu

- ▶ Recap
- ▶ Designing Servers
- ▶ Background: Registering processes
- ▶ Background: Exceptions again
- ▶ Linking processes
- ▶ Library code for making robust servers

Recap

- ▶ Organise your code in modules
- ▶ Make sure that you understand the basic concurrency primitives in Erlang
- ▶ Review the `fib` and `cfib` functions in the `fib` module.

Dealing With State

- ▶ Functions are pure (stateless).
- ▶ Processes are stateful.
- ▶ We organise our code as micro-servers that manage a state which can be manipulated via a client API.
- ▶ Functions starts processes, processes runs functions, functions are defined in modules.

Client-Server Basic Request-Response

- ▶ A server is process that loops (potentially) forever.
- ▶ Clients communicate through a given API/Protocol
- ▶ That is, we start with the following template:

```
start() -> spawn(fun () -> loop(Initial) end).  
request_reply(Pid, Request) ->  
    Pid ! {self(), Request},  
    receive  
        {Pid, Response} -> Response  
    end.  
loop(State) ->  
    receive  
        {From, Request} ->  
            {NewState, Res} = ComputeResult(Request, State),  
            From ! {self(), Res},  
            loop(NewState)  
    end.
```

Example: Position Server

```
start(Start) -> spawn(fun () -> loop(Start) end).
move(Pid, Dir) -> request_reply(Pid, {move, Dir}).
get_pos(Pid) -> request_reply(Pid, get_pos).
request_reply(Pid, Request) ->
    Pid ! {self(), Request},
    receive
        {Pid, Response} -> Response
    end.
loop({X,Y}) ->
    receive
        {From, {move, north}} ->
            From ! {self(), ok},
            loop({X, Y+1});
        {From, {move, west}} ->
            From ! {self(), ok},
            loop({X-1, Y});
        {From, {get_pos}} ->
            From ! {self(), {X,Y}},
            loop({X,Y})
    end
```

Student activity: Count Server

- ▶ Let's make a server that can keep track of a counter
- ▶ What is the client API?
- ▶ What is the internal state?

Example: Phone-Book, Interface

```
start() -> spawn(fun() -> loop(#{}) end).
```

```
add(Pid, Contact) ->  
  request_reply(Pid, {add, Contact}).
```

```
list_all(Pid) ->  
  request_reply(Pid, list_all).
```

```
update(Pid, Contact) ->  
  request_reply(Pid, {update, Contact}).
```


Example: Phone-Book, Implementation 1

```
loop(Contacts) ->
  receive
    {From, {add, Contact}} ->
      {Name, _, _} = Contact,
      case maps:is_key(Name, Contacts) of
        false ->
          From ! {self(), ok},
          loop(Contacts#{Name => Contact});
        true ->
          From ! {self(), {error, Name, is_already_there}},
          loop(Contacts)
      end;
  end;
```

Example: Phone-Book, Implementation 2

```
{From, list_all} ->
  List = maps:to_list(Contacts),
  From ! {self(),
          {ok, lists:map(fun({_, C}) -> C end, List)}},
  loop(Contacts);
{From, {update, Contact}} ->
  {Name, _, _} = Contact,
  NewContacts = maps:remove(Name, Contacts),
  From ! {self(), ok},
  loop(maps:put(Name, Contact, NewContacts));
{From, Other} ->
  From ! {self(), {error, unknow_request, Other}},
  loop(Contacts)
end.
```

Communication Patterns

- **Synchronous** (aka **Blocking**), like the simple Request-Response function blocking (aka `request_reply`).

```
blocking(Pid, Request) ->  
    Pid ! {self(), Request},  
    receive  
        {Pid, Response} -> Response  
    end.
```

- **Asynchronous** (aka **Non-Blocking**), standard sending of messages in Erlang

```
async(Pid, Msg) ->  
    Pid ! Msg.
```

Design “Method”

- ▶ Determine the API:
 - ▶ names
 - ▶ types
 - ▶ blocking or non-blocking
- ▶ Design internal protocols
- ▶ Split into servers (processes)

Design Task

You are an intern at MicroCorp. Your team has decided that you need a new logging mechanism which potentially can be distributed across machines (thus Erlang is the tool of choice). To start with, the logging mechanism should have basic functionality: it should be possible to log something (log), and it should be possible to retrieve what has been logged (retrieve).

What is your design?

Background: Registering Processes

- ▶ It can be convenient to register a process under a global name, so that we can easily get it without threading a pid around.
- ▶ `register(Name, Pid)` registers the process with `Pid` under `Name`
- ▶ `whereis(Name)` gives us the pid registered for `Name`; or undefined if `Name` is not registered
- ▶ `Name ! Mesg` sends `Mesg` to the process registered under `Name`.

Background: Exceptions Revisited

Exceptions comes in different flavours:

```
try Expr of
  Pat1 -> Expr1;
  Pat2 -> Expr2;
  ...
catch
  ExType1: ExPat1 -> ExExpr1;
  ExType2: ExPat2 -> ExExpr2;
  ...
after
  AfterExpr
end
```

Where ExType is either throw, exit, or error (throw is the default).



Background: Generating Exceptions

- ▶ We can generate all kind of exceptions:
 - ▶ `throw(Why)` to throw an exception that the caller might want to catch. For normal exceptions.
 - ▶ `exit(Why)` to exit the current process. If not caught then the message `{'EXIT',Pid,Why}` is broadcast to all *linked* processes.
 - ▶ `erlang:error(Why)` for internal errors, that nobody is really expected to handle.
- ▶ Thus, to catch *all* exceptions we need the following pattern:

```
try Expr
catch
  _ : _ -> MightyHandler
end
```


Robust Systems

- ▶ We need at least two computers(/nodes/processes) to make a robust system: one computer(/node/process) to do what we want, and one to monitor the other and take over when errors happens.
- ▶ `link(Pid)` makes a symmetric link between the calling process and `Pid`.
- ▶ `monitor(process, Pid)` makes an asymmetric link between the calling process and `Pid`.



Linking Processes

- ▶ If we want to handle when a linked process crashes then we need to call `process_flag(trap_exit, true)`.
- ▶ Thus, we have the following idioms for creating processes:

- ▶ Idiom 1, I don't care:

```
Pid = spawn(fun() -> ... end)
```

- ▶ Idiom 2, I won't live without her:

```
Pid = spawn_link(fun() -> ... end)
```

- ▶ Idiom 3, I'll handle the mess-ups:

```
process_flag(trap_exit, true),  
Pid = spawn_link(fun() -> ... end),  
loop(...).
```

```
loop(State) ->  
  receive  
    {'EXIT', Pid, Reason} -> HandleMess, loop(State);  
    ...  
  end
```

Example: Keep Trucking Looping

Suppose that we really must have a phonebook server running at all times. How do we monitor the phonebook server and restart it if (when?) it crashes.

Example: Keep Looping

```
start() -> keep_looping().
request_reply(Pid, Request) -> Ref = make_ref(),
                               Pid ! {self(), Ref, Request},
                               receive {Ref, Response} -> Response end.

keep_looping() ->
  spawn(fun () ->
    process_flag(trap_exit, true),
    Worker = spawn_link(fun() -> loop({}) end),
    supervisor(Worker)
  end).

supervisor(Worker) ->
  receive
    {'EXIT', Pid, Reason} ->
      io:format("~p exited because of ~p~n", [Pid, Reason]),
      Pid1 = spawn_link(fun() -> loop({}) end),
      supervisor(Pid1);
    Msg -> Pid ! Msg, supervisor(Pid)
  end.
```

Distributed Programs in Erlang

- ▶ *Distributed Erlang* for tightly coupled computers in a secure environment.
 - ▶ `spawn(Node, Fun)` to spawn a process running `Fun` on `Node`
 - ▶ `{RegAtom, Node} ! Mesg` sends `Mesg` to the process registered as `RegAtom` at `Node`.
 - ▶ `monitor_node(Node, Flag)` register the calling process for notification about `Node` if `Flag` is `true`; if `Flag` is `false` then monitoring is turned off.
- ▶ *Sockets* for untrusted environments:
 - ▶ To build a middleware layer for Erlang nodes
 - ▶ For inter-language communication.

See the documentation for `gen_tcp` and `gen_udp`

Setting Up Some Erlang Nodes

- ▶ To start nodes on the same machine, start `erl` with option `-sname`
- ▶ To start nodes on different machines, start `erl` with options `-name` and `-setcookie`:
 - ▶ On machine A:
`erl -name bart -setcookie BoomBoomShakeTheRoom`
 - ▶ On machine B:
`erl -name homer -setcookie BoomBoomShakeTheRoom`
- ▶ `rpc:call(Node, Mod, Fun, Args)` evaluates `Mod:Fun(Args)` on `Node`. (See the the manual page for `rpc` for more information.)

Generic Servers

- ▶ Goal: Abstract out the difficult handling of concurrency to a generic library
- ▶ The difficult parts:
 - ▶ The `start-request_reply/async-loop` pattern
 - ▶ Supervisors

Basic Server Library

```
start(Name, Mod) ->
  register(Name, spawn(fun() -> loop(Name, Mod, Mod:init())
                end)).
request_reply(Pid, Request) ->
  Pid ! {self(), Request},
  receive
    {Pid, Reply} -> Reply
  end.
loop(Name, Mod, State) ->
  receive
    {From, Request} ->
      {Reply, State1} = Mod:handle(Request, State),
      From ! {Name, Reply},
      loop(Name, Mod, State1)
  end.
```


Example: Phonebook Callback Module, 1

```
-module(pb).  
-import(basicserver, [request_reply/2]).
```

%% Interface

```
start()           -> basicserver(phonebook, pb).  
add(Contact)     -> request_reply(pb, {add, Contact}).  
list_all()        -> request_reply(pb, list_all).  
update(Contact)  -> request_reply(pb, {update, Contact}).
```

Example: Phonebook Callback Module, 2

%% Callback functions

```
init() -> #{}
```

```
handle({add, {Name, _, _} = Contact}, Contacts) ->  
  case maps:is_key(Name, Contacts) of  
    false -> {ok, Contacts#{Name => Contact}};  
    true -> {{error, Name, is_already_there},  
            Contacts}
```

```
end;
```

```
handle(list_all, Contacts) ->  
  List = maps:to_list(Contacts),  
  {{ok, lists:map(fun(_, C) -> C end, List)},  
   Contacts};
```

```
handle({update, {Name, _, _} = Contact}, Contacts) ->  
  {ok, Contacts#{Name => Contact}}.
```

Say something about
Kaboose!

Summary

- ▶ How design micro-server: blocking vs non-blocking
- ▶ To make a robust system we need two parts: one to do the job and one to take over in case of errors
- ▶ Structure your code into the infrastructure parts and the functional parts.